

Device

```
=== Device Information ===
Available Devices:
Device 0: "NVIDIA GeForce RTX 4050 Laptop GPU"
Selected Device: NVIDIA GeForce RTX 4050 Laptop GPU
Selected Device ID: 0
Compute Capability: 8.9
SMs: 20
Device Global Memory: 6140 MiB
Shared Memory per SM: 100 KiB
Memory Bus Width: 96 bits (ECC disabled)
Application Compute Clock Rate: 2.355 GHz
Application Memory Clock Rate: 7.825 GHz
```

Output Details

- **Latency**: refers to the [min, max, mean, median, 99% percentile] of the engine latency measurements, when timing the engine w/o profiling layers.
- **Throughput**: is measured in query (inference) per second (QPS).
- **Enqueue Time**: Time taken to enqueue inference requests.
- **H2D Latency**: Host-to-Device latency (data transfer time).
- **GPU Compute Time**: Time spent computing on the GPU.
- **D2H Latency**: Device-to-Host latency (data transfer time).
- **Total Host Walltime**
- **Total GPU Compute Time**

YOLOv9-C QAT (Mish)

Precision: FP32+FP16+INT8

Batch Size 1

```
=== Performance summary ===
Throughput: 347.915 qps
```

```
Latency: min = 2.85596 ms, max = 2.89282 ms, mean = 2.87257 ms, median = 2.87305 ms, percentile(90%) = 2.87524 ms, percentile(95%) = 2.87598 ms, percentile(99%) = 2.87744 ms
```

```
Enqueue Time: min = 0.0185547 ms, max = 2.50635 ms, mean = 0.0458279 ms, median = 0.0307617 ms, percentile(90%) = 0.079834 ms, percentile(95%) = 0.0966797 ms, percentile(99%) = 0.199219 ms
```

```
H2D Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
GPU Compute Time: min = 2.85596 ms, max = 2.89282 ms, mean = 2.87257 ms, median = 2.87305 ms, percentile(90%) = 2.87524 ms, percentile(95%) = 2.87598 ms, percentile(99%) = 2.87744 ms
```

```
D2H Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
Total Host Walltime: 10.0082 s
```

```
Total GPU Compute Time: 10.0023 s
```

BatchSize 4

```
=== Performance summary ===
```

```
Throughput: 100.77 qps
```

```
Latency: min = 9.89899 ms, max = 9.94409 ms, mean = 9.92196 ms, median = 9.92188 ms, percentile(90%) = 9.92871 ms, percentile(95%) = 9.92969 ms, percentile(99%) = 9.93506 ms
```

```
Enqueue Time: min = 0.0205078 ms, max = 0.917236 ms, mean = 0.0534803 ms, median = 0.0455322 ms, percentile(90%) = 0.0944824 ms, percentile(95%) = 0.106689 ms, percentile(99%) = 0.139404 ms
```

```
H2D Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
GPU Compute Time: min = 9.89899 ms, max = 9.94409 ms, mean = 9.92196 ms, median = 9.92188 ms, percentile(90%) = 9.92871 ms, percentile(95%) = 9.92969 ms, percentile(99%) = 9.93506 ms
```

```
D2H Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
Total Host Walltime: 10.0229 s
```

```
Total GPU Compute Time: 10.0212 s
```

BatchSize 8

```
=== Performance summary ===
```

```
Throughput: 48.9507 qps
```

```
Latency: min = 20.3643 ms, max = 21.1794 ms, mean = 20.4271 ms, median = 20.4258 ms, percentile(90%) = 20.437 ms, percentile(95%) = 20.4409 ms, percentile(99%) = 20.4473 ms
```

```
Enqueue Time: min = 0.0214844 ms, max = 0.42627 ms, mean = 0.075538 ms, median =
```

```
0.0620117 ms, percentile(90%) = 0.117676 ms, percentile(95%) = 0.127441 ms,  
percentile(99%) = 0.233887 ms
```

```
H2D Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) =  
0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
GPU Compute Time: min = 20.3643 ms, max = 21.1794 ms, mean = 20.4271 ms, median =  
20.4258 ms, percentile(90%) = 20.437 ms, percentile(95%) = 20.4409 ms,  
percentile(99%) = 20.4473 ms
```

```
D2H Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) =  
0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
Total Host Walltime: 10.0305 s
```

```
Total GPU Compute Time: 10.0297 s
```

BatchSize 12

```
=== Performance summary ===
```

```
Throughput: 32.2187 qps
```

```
Latency: min = 30.9883 ms, max = 31.0742 ms, mean = 31.0362 ms, median = 31.0371  
ms, percentile(90%) = 31.0498 ms, percentile(95%) = 31.0537 ms, percentile(99%) =  
31.063 ms
```

```
Enqueue Time: min = 0.0253906 ms, max = 0.40625 ms, mean = 0.0721937 ms, median =  
0.0598145 ms, percentile(90%) = 0.114075 ms, percentile(95%) = 0.131348 ms,  
percentile(99%) = 0.256836 ms
```

```
H2D Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) =  
0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
GPU Compute Time: min = 30.9883 ms, max = 31.0742 ms, mean = 31.0362 ms, median =  
31.0371 ms, percentile(90%) = 31.0498 ms, percentile(95%) = 31.0537 ms,  
percentile(99%) = 31.063 ms
```

```
D2H Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) =  
0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
Total Host Walltime: 10.0563 s
```

```
Total GPU Compute Time: 10.0557 s
```

BatchSize 16

```
=== Performance summary ===
```

```
Throughput: 24.0033 qps
```

```
Latency: min = 41.6174 ms, max = 41.7002 ms, mean = 41.6591 ms, median = 41.6594 ms, percentile(90%) = 41.6748 ms, percentile(95%) = 41.679 ms, percentile(99%) = 41.6892 ms
```

```
Enqueue Time: min = 0.0283203 ms, max = 1.18127 ms, mean = 0.0669195 ms, median = 0.057312 ms, percentile(90%) = 0.0908203 ms, percentile(95%) = 0.109375 ms, percentile(99%) = 0.132812 ms
```

```
H2D Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
GPU Compute Time: min = 41.6174 ms, max = 41.7002 ms, mean = 41.6591 ms, median = 41.6594 ms, percentile(90%) = 41.6748 ms, percentile(95%) = 41.679 ms, percentile(99%) = 41.6892 ms
```

```
D2H Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
Total Host Walltime: 10.0819 s
```

```
Total GPU Compute Time: 10.0815 s
```

BatchSize Resilience

```
=== Performance summary ===
```

```
Throughput: 48.8766 qps
```

```
Latency: min = 20.4316 ms, max = 20.4822 ms, mean = 20.4578 ms, median = 20.458 ms, percentile(90%) = 20.4697 ms, percentile(95%) = 20.4736 ms, percentile(99%) = 20.4795 ms
```

```
Enqueue Time: min = 0.0227051 ms, max = 0.307617 ms, mean = 0.0747259 ms, median = 0.0644531 ms, percentile(90%) = 0.120117 ms, percentile(95%) = 0.136719 ms, percentile(99%) = 0.188965 ms
```

```
H2D Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
GPU Compute Time: min = 20.4316 ms, max = 20.4822 ms, mean = 20.4578 ms, median = 20.458 ms, percentile(90%) = 20.4697 ms, percentile(95%) = 20.4736 ms, percentile(99%) = 20.4795 ms
```

```
D2H Latency: min = 0 ms, max = 0 ms, mean = 0 ms, median = 0 ms, percentile(90%) = 0 ms, percentile(95%) = 0 ms, percentile(99%) = 0 ms
```

```
Total Host Walltime: 10.0457 s
```

```
Total GPU Compute Time: 10.0448 s
```